Lesson 7: Linear Regression Model – Part 1

◇ Introduction

Welcome to one of the most fundamental and widely used models in machine learning — Linear Regression. Linear Regression isn't just a model — it's a powerful concept that lays the groundwork for many advanced techniques. Whether you're predicting house prices, forecasting sales, or analyzing trends, linear regression is often the first model you'll reach for.
In this lesson, we'll introduce the concept of linear regression, look at practical examples, and break down key terminology to build a solid foundation.

◇ Examples: Where Do We Use Linear Regression?

- 🏠 House Price Prediction
- Predict the price of a house based on square footage.
  ➡️ More area → Higher price (usually a linear relationship).
- 📈 Sales Forecasting
- Estimate future sales based on advertising spend or seasonal trends.
  ➡️ More ads → More sales.
- 🌡️ Temperature Prediction
  Forecast tomorrow's temperature based on historical data.

These are just a few examples — linear regression can be used anywhere there's a linear relationship between input and output.

◇ What is Linear Regression?

At its core, linear regression is a supervised learning algorithm used for predicting a continuous output.
It tries to find the best-fitting straight line that describes the relationship between the independent variable(s) (features) and the dependent variable (target).

◇ The Linear Regression Equation

In simple linear regression (1 feature):

$$y=w\cdot x+b$$

- x: Input feature (e.g., square footage)
- y: Predicted output (e.g., price)
- w: Weight (slope of the line)
- b: Bias (y-intercept)

The goal is to find values of w and b that minimize the error between the predicted y and the actual y.

◇ Key Terminology

Let’s go over some important terms you’ll hear often when working with linear regression:

Term	Description
Feature (x)	The independent variable(s), also called input(s)
Target (y)	The dependent variable, the value we are trying to predict
Weight (w)	The slope — how much y changes for a change in x
Bias (b)	The y-intercept — where the line crosses the y-axis
Loss Function	Measures how far predictions are from actual values
Cost Function	The average of the loss across all examples (often MSE: Mean Squared Error)
Gradient Descent	The optimization method used to find the best w and b

---

◇ Visualizing Linear Regression

Imagine plotting data points on a 2D graph — Linear Regression tries to draw a line those best fits all those points.

If the relationship is strong and linear, most points will be close to the line. If not, the line might not fit well — and that's where more advanced models come in.

◇ Outro

In this lesson, we covered:

- The basics of what linear regression is and why it's important.
- Real-world examples where it's used.
- The mathematical equation and terminology behind the model.

In Part 2, we’ll go deeper — you’ll learn how to train a linear regression model, explore evaluation metrics, and implement it in code using tools like Python and Scikit-learn.